

MAT 170 Homework Project 3

Joe Lenning Student id: 919484830

May 19, 2024

1 Problem 3A

1.1 Code for Matrix

```
1 import csv
2 import datetime
3 import numpy as np
4 import os
5
6 def read_values_hist(RAFI_combined):
7     with open('RAFI_combined', 'r') as my_file:
8         reader = csv.DictReader(RAFI_combined)
9         mydict = {}
10        for row in reader:
11            # Convert the date format from the CSV and capture the values
12            date = datetime.datetime.strptime(row['Date'], '%m/%d/%Y')
13            mydict[date] = float(row['Value_With_Dividends__USD'])
14        return mydict
15
16 def next_month(d):
17     return datetime.datetime(d.year + (d.month // 12), ((d.month % 12) + 1), 1)
18
19 def find_closest_date(data, target_date):
20     """Find the closest date within four days before or after the 1st of the month."""
21     for offset in range(-4, 5): # From four days before to four days after
22         close_date = target_date + datetime.timedelta(days=offset)
23         if close_date in data:
24             return close_date
25     return None
```

```

26
27 # Paths to the data files
28 file_paths = {
29     "RAFI Global": 'path/to/RAFI Global.csv',
30     "RAFI US": 'path/to/RAFI US.csv',
31     "RAFI EU": 'path/to/RAFI EU.csv',
32     "RAFI Emerging Markets": 'path/to/RAFI Emerging Markets.csv',
33     "RAFI Developed ex-US": 'path/to/RAFI Developed US.csv'
34 }
35
36 #data from files
37 index_data = {name: read_values_hist(path) for name, path in file_paths.items()}
38
39 #dates
40 start_date = datetime.datetime(1996, 7, 1)
41 end_date = datetime.datetime(2023, 12, 1)
42
43 all_dates = [start_date + datetime.timedelta(days=(next_month(start_date) - start_date).days * i) for i in range
44
45 returns_matrix = np.zeros((len(file_paths), len(all_dates) - 1)) # 5 indices and 329 months
46
47 for i, (name, data) in enumerate(index_data.items()):
48     for j in range(1, len(all_dates)):
49         prev_date = find_closest_date(data, all_dates[j-1])
50         curr_date = find_closest_date(data, all_dates[j])
51         if prev_date and curr_date:
52             # Calculate return as (V(t) / V(t_prev)) - 1
53             returns_matrix[i, j-1] = (data[curr_date] / data[prev_date]) - 1
54
55 # Center the matrix by subtracting the mean of each row
56 returns_centered = returns_matrix - np.mean(returns_matrix, axis=1, keepdims=True)
57
58 # Outputs the full 5x329 matrix centered by the mean
59 returns_centered

```

1.2 PCA Model

First we compute the covariance matrix Σ of the data X :

$$\Sigma = \frac{1}{n-1} X^T X$$

The first principal component vector v_1 is obtained by solving the following optimization problem:

$$\max_v v^T \Sigma v \quad \text{subject to} \quad \|v\| = 1$$

This optimization problem maximizes the variance of the projected data and is subject to the constraint that v has unit length.

The principal vector v_1 corresponds to the eigenvector of Σ associated with the largest eigenvalue. This can be derived from the eigenvalue equation:

$$\Sigma v = \lambda v$$

where λ is an eigenvalue of Σ and v is the corresponding eigenvector.

1.3 U Σ for Data

The following Python code uses the `numpy` library to perform Singular Value Decomposition (SVD) and extracts the matrices U , Σ , and V^T :

```
1 import numpy as np
2
3 # Assuming 'returns_centered' is the centered matrix from your data processing
4 U, Sigma, VT = np.linalg.svd(returns_centered, full_matrices=False)
```

The matrix U and diagonal values from Σ are obtained as outputs from the SVD operation.

$$U = \begin{bmatrix} -0.4013 & -0.2745 & 0.2612 & 0.1065 & -0.8271 \\ -0.3409 & -0.4085 & 0.4009 & -0.6626 & 0.3423 \\ -0.4894 & -0.1078 & -0.8417 & -0.2005 & -0.0184 \\ -0.5499 & 0.8005 & 0.2252 & -0.0397 & 0.0672 \\ -0.4253 & -0.3245 & 0.1095 & 0.7126 & 0.4404 \\ \vdots & \vdots & \vdots & \vdots & \vdots \end{bmatrix}$$

$$\Sigma = \begin{pmatrix} 1.2009 & 0 & 0 & 0 & 0 \\ 0 & 0.4053 & 0 & 0 & 0 \\ 0 & 0 & 0.3047 & 0 & 0 \\ 0 & 0 & 0 & 0.2097 & 0 \\ 0 & 0 & 0 & 0 & 0.0173 \\ \vdots & \vdots & \vdots & \vdots & \vdots \end{pmatrix}$$

1.4 Error Term

Again we can just add the following snippet to our code

```
1 total_variance = np.sum(sigma**2)
2 eta_k = (sigma**2) / total_variance
3
4 # Output the eta_k ratios
5 print("Ratios _k:", eta_k)
```

We see the result as The ratios of explained variance for each principal component derived from the singular values are:

$$\begin{aligned}\eta_1 &\approx 0.8271 \quad (82.71\%) \\ \eta_2 &\approx 0.0942 \quad (9.42\%) \\ \eta_3 &\approx 0.0532 \quad (5.32\%) \\ \eta_4 &\approx 0.0252 \quad (2.52\%) \\ \eta_5 &\approx 0.0002 \quad (0.02\%)\end{aligned}$$

These values indicate how much variance 1 component takes from the total variance of the data set. A high η_k suggests that a significant portion of the data's variation can be observed and analyzed effectively when projected onto the corresponding k -dimensional subspace(per the book). This insight is particularly useful for dimensionality reduction, as it helps to identify how many components should be retained to effectively capture the main characteristics while reducing rank.

2 Problem 3B, Least squares via pseudoinverse

2.1 a

The set of all optimal solutions to the least squares problem $\min \|Ax - b\|^2$ can be denoted as χ , where:

$$\chi = \{x \mid x = x^* + v, \text{ for all } v \in \text{Null}(A)\}$$

Here, x^* is any particular solution to the equation $Ax = \text{Proj}_{\text{Range}(A^T)}(b)$ and $\text{Null}(A)$ represents the null space of A .

2.2 b

Let $A \in R^{m \times n}$ and assume A has full rank. The solution to the least squares problem, $x^* = A^+b$, can be shown to be optimal by the following derivation:

First, consider the least squares problem formulated as minimizing $\|Ax - b\|^2$. For the solution $x^* = A^+b$, where A^+ is the Moore-Penrose pseudoinverse of A , we calculate:

$$\|Ax^* - b\|^2 = \|A(A^+b) - b\|^2 = \|AA^+b - b\|^2 = \|Pb - b\|^2 = \|0\|^2 = 0,$$

where $P = AA^+$ is the orthogonal projection matrix onto the range of A . The operation AA^+b simplifies to Pb because AA^+ acts as the projection onto the range of A , simplifying the expression further to Pb . Since $Pb = b$ for all b in the column space of A , it implies that the residuals $Pb - b = 0$, affirming that x^* indeed minimizes the squared error to zero.

Additionally we can verify this through the gradient condition. For the function $f(x) = \|Ax - b\|_2^2$, the optimal solution must satisfy:

$$\nabla f(x) = 0.$$

The gradient of $f(x)$ with respect to x is given by:

$$\nabla f(x) = 2A^T(Ax - b).$$

Setting this to zero for the optimal condition, we have:

$$2A^T(Ax - b) = 0 \Rightarrow A^T(Ax - b) = 0.$$

Substituting $x = x^*$:

$$A^T(Ax^* - b) = A^T(AA^+b - b) = A^T(Pb - b) = A^T(0) = 0,$$

which confirms that the gradient at x^* is zero, therefore, satisfying the necessary condition for a minimum.

Thus, $x^* = A^+b$ is not only a solution that minimizes $\|Ax - b\|^2$ but also satisfies the gradient condition for optimality in an unconstrained optimization framework. This dual verification through both projection and gradient analysis robustly establishes x^* as the optimal solution.

2.3 c

Let $A \in R^{m \times n}$ and $b \in R^m$. Consider the least squares problem:

$$\min_x \|Ax - b\|^2.$$

Given a matrix $A \in R^{m \times n}$ and a vector $b \in R^m$, we aim to solve the least squares problem by finding $x^* = A^+b$, where A^+ denotes the pseudoinverse of A .

The solution $x^* = A^+b$ is optimal for the problem:

$$\min_x \|Ax - b\|^2.$$

To verify the optimality and properties of x^* , we consider two key conditions:

1. $d \in \text{Range}(C)$
2. $C^T d = A^T b$

where $d = A^T b$ and $C = A^T A$.

Using the singular value decomposition of A , express A as $A = U\Sigma V^T$, where U and V are orthogonal matrices, and Σ is a diagonal matrix of singular values.

Thus, C can be expressed as:

$$C = A^T A = V\Sigma^T \Sigma V^T.$$

And the vector d is:

$$d = A^T b = V \Sigma^T U^T b.$$

Since V spans the range of C , and d is expressed in terms of V and Σ , it follows that d is in the range of C , confirming that:

$$d \in \text{Range}(C).$$

Given that the columns of V span the range of C , d is clearly in the range of C , since:

$$\text{Range}(C) = \text{Range}(V).$$

Thus the solution $x^* = A^+ b$ is a solution to the one with the minimum 2-norm among all solutions, proving it to be the optimal solution in the sense of the 2-norm minimization.

2.4 Extra Credit

Prove that if $d \in \text{Range}(C)$ (i.e., $Cx = d$ is feasible), then $x^* = C^+ d$ is an optimal solution to the 2-norm minimization problem $\min \|x\|_2$ subject to $Cx = d$:

We prove that if d is in the range of matrix C , and the equation $Cx = d$ is feasible, then the solution $x^* = C^+ d$ provided by the pseudoinverse C^+ is the optimal solution to the 2-norm minimization problem:

$$\min \|x\|_2 \quad \text{subject to} \quad Cx = d.$$

The pseudoinverse C^+ of a matrix C has several crucial properties, notably that $C^+ C$ and $C C^+$ are projection matrices. These properties facilitate the optimal solution as follows:

- $C^+ C$ projects onto the range (column space) of C^T .
- $C C^+$ projects onto the range of C .

The operation $x^* = C^+ d$ projects d onto the column space of C^T . This is significant because it ensures that x^* lies in the intersection of the column space of C^T and the affine space defined by the equation $Cx = d$, and it is the vector in this intersection with the minimal 2-norm.

The projection can be expressed mathematically as:

$$x^* = C^+ d = C^+ C x^*.$$

This expression confirms that x^* is not only in the column space of C^T but is also the image under the projection matrix $C^+ C$, which projects vectors onto this space. This solution is the optimal 2-norm solution because it minimizes the Euclidean distance to the origin while satisfying $Cx = d$.

The pseudoinverse solution $x^* = C^+ d$ is shown to be optimal in the sense that it minimizes the 2-norm among all solutions that satisfy $Cx = d$. The proof hinges on the projection properties of C^+ and the orthogonality principle, which ensure that x^* has no components in the direction of any vectors that lie in the null space of C .

3 Problem 3C

3.1 Mathematical Expression

Given a dataset $\{(x_i, y_i)\} \subseteq R^2$ where $i = 1, 2, \dots, m$, we aim to fit a polynomial model of degree k . The polynomial model is defined by the function:

$$y = \alpha^T \phi(x)$$

where $\Phi = [\phi_0, \phi_1, \dots, \phi_k]^T$ is the vector of coefficients, and $x^{(k)} = [1, x, x^2, \dots, x^k]^T$ represents the powers of x .

To express this in a matrix form suitable for least squares regression, we use our Φ above however this time in matrix form:

$$\Phi = \begin{bmatrix} x_1^k & x_1^{k-1} & \dots & x_1 & 1 \\ x_2^k & x_2^{k-1} & \dots & x_2 & 1 \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ x_m^k & x_m^{k-1} & \dots & x_m & 1 \end{bmatrix}$$

This matrix representation allows us to model the polynomial relationship and apply least squares regression to estimate the polynomial coefficients by minimizing the squared residuals between the predicted values and the observed data.

3.2 Mathematical Expression

Generally, our regularization term is as follows:

$$\min_{\alpha} \|\Phi\alpha - y\|_2^2 + \mathcal{L}(\alpha)$$

And if we want to apply L2 norm coefficients in the polynomial model, we get:

$$\min \|\Phi\alpha - y\|_2^2 + \lambda \|\alpha\|_2^2$$

And the closed form solution:

$$\Phi^* = (X^T X + \lambda I)^{-1} X^T y$$

3.3 Code

```
1 import numpy as np
2 from sklearn.model_selection import train_test_split
3 import matplotlib.pyplot as plt
4 import cvxpy as cp
5
6 # Generate the data set, containing 100 instances
7 np.random.seed(seed=0)
8 data_size = 100
```

```

9 data_x = np.linspace(0, 2*np.pi, num=data_size)
10 data_y = np.apply_along_axis(lambda x: np.sin(x), 0, data_x) + 0.25 * np.random.normal(size=data_x.shape)
11
12 # A figure of the underlying generator function and the data points
13 plt.scatter(data_x, data_y, color="black")
14 xlist = np.linspace(0, 2*np.pi, num=100)
15 plt.plot(xlist, np.apply_along_axis(lambda x: np.sin(x), 0, xlist), color="green")
16
17 # Splitting the training and test set
18 X_train, X_test, y_train, y_test = train_test_split(data_x, data_y, test_size=0.2, random_state=0)
19
20 # Define the degree of the polynomial
21 k = 3
22
23 # Create a Vandermonde matrix for X_train
24 n = len(X_train)
25 X_vander = np.vander(X_train, N=k+1, increasing=True)
26
27 # Define the CVXPY variables for polynomial coefficients
28 coeffs = cp.Variable(k+1)
29
30 # Define the objective function (least squares)
31 objective = cp.Minimize(cp.sum_squares(X_vander @ coeffs - y_train))
32
33 # Define the problem and solve
34 problem = cp.Problem(objective)
35 problem.solve(solver=cp.SCS)
36
37 # Output the learned polynomial coefficients
38 print("The polynomial coefficients are:", coeffs.value)
39
40 # Generate a dense grid of x values for plotting the fitted polynomial
41 x_dense = np.linspace(0, 2*np.pi, 400)
42 X_vander_dense = np.vander(x_dense, N=k+1, increasing=True)
43
44 # Evaluate the polynomial at the grid points
45 y_dense = X_vander_dense @ coeffs.value
46
47 plt.scatter(X_train, y_train, color="black", label="Data points")
48 plt.plot(x_dense, y_dense, label=f"Fitted Polynomial (degree {k})", color="blue")
49 plt.legend()
50 plt.show()

```


3.4 Code

```
1 import numpy as np
2 from sklearn.model_selection import train_test_split
3 import matplotlib.pyplot as plt
4 import cvxpy as cp
5
6 # Generate the data set, containing 100 instances
7 np.random.seed(seed=0)
8 data_size = 100
9 data_x = np.linspace(0, 2*np.pi, num=data_size)
10 data_y = np.apply_along_axis(lambda x: np.sin(x), 0, data_x) + 0.25 * np.random.normal(size=data_x.shape)
11
12 # A figure of the underlying generator function and the data points
13 plt.scatter(data_x, data_y, color="black")
14 xlist = np.linspace(0, 2*np.pi, num=100)
15 plt.plot(xlist, np.apply_along_axis(lambda x: np.sin(x), 0, xlist), color="green")
16
17 # Splitting the training and test set
18 X_train, X_test, y_train, y_test = train_test_split(data_x, data_y, test_size=0.2, random_state=0)
19
20 # Define the degree of the polynomial
21 k = 12 # Degree of the polynomial
22
23 # Create a Vandermonde matrix for X_train
24 X_vander = np.vander(X_train, N=k+1, increasing=True)
25
26 # Define the CVXPY variables for polynomial coefficients
27 coeffs = cp.Variable(k+1)
28
29 # Define the regularization parameter
30 lambda_reg = 1
31
32 # Define the objective function (least squares + L1 regularization)
33 objective = cp.Minimize(cp.sum_squares(X_vander @ coeffs - y_train) + lambda_reg * cp.norm1(coeffs))
34
35 # Define the problem and solve
36 problem = cp.Problem(objective)
37 problem.solve(solver=cp.SCS)
38
39 # Output the learned polynomial coefficients
40 print("The polynomial coefficients are:", coeffs.value)
41
42 # Generate a dense grid of x values for plotting the fitted polynomial
43 x_dense = np.linspace(0, 2*np.pi, 400)
44 X_vander_dense = np.vander(x_dense, N=k+1, increasing=True)
45
```

```

46 # Evaluate the polynomial at the grid points
47 y_dense = X_vander_dense @ coeffs.value
48
49 plt.scatter(X_train, y_train, color="black", label="Data points")
50 plt.plot(x_dense, y_dense, label=f"Fitted Polynomial (degree {k})", color="blue")
51 plt.legend()
52 plt.show()
53
54 # Calculate and print the test error
55 X_vander_test = np.vander(X_test, N=k+1, increasing=True)
56 test_error = np.mean((X_vander_test @ coeffs.value - y_test) ** 2)
57 print("Test error:", test_error)
58
59 # Check the sparsity of the optimal solution
60 sparsity = np.sum(coeffs.value == 0)
61 print("Number of zero coefficients (sparsity):", sparsity)

```

3.5 Extra credit

NA

4 Problem 3D

4.1 Mathematical Expression Code

The mathematical expression for the minimization problem in the Support Vector Machine (SVM) model is given by:

$$\min_{w, \beta} \frac{1}{2} \|w\|^2$$

subject to the constraint that for all training samples,

$$y_i(w^T x_i + \beta) \geq 1, \quad \forall i = 1, 2, \dots, m$$

Here, w is the weight vector and β (often denoted as b and misread as η) is the bias term of the hyperplane. The vectors x_i represent the feature vectors of the training samples, and y_i are the labels of these samples, which are typically $+1$ or -1 , corresponding to the two classes.

```

1 import numpy as np
2 from sklearn import datasets
3 from sklearn.model_selection import train_test_split
4 import matplotlib.pyplot as plt
5 solver=cp.ECOS

```

```

6 WDBC = datasets.load_breast_cancer()
7 X, y = datasets.load_breast_cancer(return_X_y=True)
8 # Changing labels
9 y = 2*y - 1
10 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=0)
11
12 # Normalize the features on the training set to have mean 0 and standard deviation 1
13 for i in range(X_train.shape[1]):
14     X_train_mean = np.mean(X_train[:,i])
15     X_train_std = np.std(X_train[:,i])
16     X_train[:,i] = (X_train[:,i] - X_train_mean)/X_train_std
17     X_test[:,i] = (X_test[:,i] - X_train_mean)/X_train_std
18
19
20 # 2-dimensional example
21 X2D_train = X_train[:, :4]
22 #X2D_test = X_test[:, :40]
23 X2D_test = X_test[:, :4]
24
25 fig, ax = plt.subplots()
26 ax.scatter(X2D_train[:,0], X2D_train[:,1], c=y_train, cmap=plt.cm.coolwarm, s=20, edgecolors="k")
27 plt.show()
28
29
30 import cvxpy as cp
31
32 # Define the variables
33 d = X2D_train.shape[1]
34 w = cp.Variable(d)
35 beta = cp.Variable()
36
37 # Set up the problem
38 objective = cp.Minimize(0.5 * cp.norm(w, 2)**2)
39 constraints = [y_train * (X2D_train @ w + beta) >= 1]
40 prob = cp.Problem(objective, constraints)
41
42 # Solve the problem
43 prob.solve()
44
45 # Print results
46 w_value = w.value
47 beta_value = beta.value
48 print("Weight vector (w):", w_value)
49 print("Bias (beta):", beta_value)
50
51 # Predicting test labels
52 predictions = X2D_test @ w_value + beta_value

```

```

53 predictions = np.where(predictions >= 0, 1, -1)
54
55 # Calculate test error
56 test_error = np.mean(predictions != y_test)
57 print("Test error:", test_error)
58 Weight vector (w): [4.21618501e-10 2.34296821e-10 4.28819818e-10 4.06486254e-10]
59 Bias (beta): 0.10956616072203625
60 Test error: 0.3706293706293706]

```

4.2 Mathematical Expression and Code

Mathematical Expression Let:

- x_i be the input features for each data point i ,
- y_i be the corresponding labels, which are ± 1 ,
- w be the weight vector of the hyperplane,
- β (often denoted as b) be the bias term,
- ξ_i be the slack variables representing the degree of misclassification for each data point i .

$$\min_{w, \beta, \xi} \left(\frac{1}{2} \|w\|^2 + C \sum_{i=1}^n \xi_i \right) \quad (1)$$

s.t

$$y_i(w^T x_i + \beta) \geq 1 - \xi_i \quad \text{for all } i \quad (2)$$

and

$$\xi_i \geq 0 \quad \text{for all } i \quad (3)$$

where C is a regularization parameter that controls the trade-off between achieving a low error on the training data and maintaining a small norm for w .

```

1  import numpy as np
2  import cvxpy as cp
3  from sklearn import datasets
4  from sklearn.model_selection import train_test_split
5  import matplotlib.pyplot as plt
6
7  #Data Prepration
8  WDBC = datasets.load_breast_cancer()
9  X, y = datasets.load_breast_cancer(return_X_y=True)
10 # Change labels from {0, 1} to {-1, 1}
11 y = 2 * y - 1

```

```

12 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=0)
13
14 for i in range(X_train.shape[1]):
15     X_train_mean = np.mean(X_train[:, i])
16     X_train_std = np.std(X_train[:, i])
17     X_train[:, i] = (X_train[:, i] - X_train_mean) / X_train_std
18     X_test[:, i] = (X_test[:, i] - X_train_mean) / X_train_std
19
20 fig, ax = plt.subplots()
21 ax.scatter(X_train[:, 0], X_train[:, 1], c=y_train, cmap=plt.cm.coolwarm, s=20, edgecolors="k")
22 plt.show()
23
24 # Setup SVM model
25 d = X_train.shape[1]
26 w = cp.Variable(d)
27 beta = cp.Variable()
28 e = cp.Variable(X_train.shape[0])
29 lambda_param = 1
30
31 # Objective and constraints
32 objective = cp.Minimize(cp.norm(e, 1) + lambda_param * cp.norm(w, 2)**2)
33 constraints = [y_train * (X_train @ w + beta) + e >= 1, e >= 0]
34 prob = cp.Problem(objective, constraints)
35
36 prob.solve(solver=cp.ECOS)
37
38 print("Weight vector (w):", w.value)
39 print("Bias (beta):", beta.value)
40
41 predictions = X_test @ w.value + beta.value
42 predictions = np.where(predictions >= 0, 1, -1)
43
44 test_error = np.mean(predictions != y_test)
45 print("Test error:", test_error)
46 Weight vector (w): [ 3.38538449e-10  1.88052705e-10  3.44325806e-10  3.26374870e-10
47  1.75975354e-10  2.86678687e-10  3.22528516e-10  3.64615422e-10
48  1.61129917e-10 -1.09939421e-11  2.61106206e-10 -1.54766486e-11
49  2.55501786e-10  2.46934770e-10 -4.21915318e-11  1.24379553e-10
50  8.76699178e-11  1.75935092e-10 -1.39383048e-11  2.51790383e-11
51  3.61229503e-10  2.08667935e-10  3.63913517e-10  3.39106418e-10
52  2.02980273e-10  2.83722806e-10  3.09453603e-10  3.73915076e-10
53  1.99963993e-10  1.51908233e-10]
54 Bias (beta): 0.07291137243980425
55 Test error: 0.3706293706293706

```