

MAT 170 Homework Project 4

Joe Lenning Student id: 919484830

June 6, 2024

1 Problem 1

1.1 Model

We our original model given by:

Maximize $1.5x + 1.75y + 2.0z$

$$\begin{aligned}9x + 4y + z &\leq 1600, \\5x + 8y + 5z &\leq 1920, \\2x + 4y + 10z &\leq 1760, \\x, y, z &\geq 0.\end{aligned}$$

However we see that our dual version is

Minimize $W = 1600y_1 + 1920y_2 + 1760y_3$

Subject to:

$$\begin{aligned}9y_1 + 5y_2 + 2y_3 &\geq 1.5, \\4y_1 + 8y_2 + 4y_3 &\geq 1.75, \\y_1 + 5y_2 + 10y_3 &\geq 2, \\y_1, y_2, y_3 &\geq 0.\end{aligned}$$

Alternatively we can:

Minimize $W = -1600y_1 - 1920y_2 - 1760y_3$

Subject to:

$$\begin{aligned}-9y_1 - 5y_2 - 2y_3 &\leq -1.5, \\ -4y_1 - 8y_2 - 4y_3 &\leq -1.75, \\ -y_1 - 5y_2 - 10y_3 &\leq -2, \\ y_1, y_2, y_3 &\leq 0.\end{aligned}$$

1.2 Code

```
1 import numpy as np
2 import cvxpy as cp
3
4 # Define the variables for the primal problem
5 x = cp.Variable()
6 y = cp.Variable()
7 z = cp.Variable()
8
9 # Define the variables for the dual problem
10 y1 = cp.Variable()
11 y2 = cp.Variable()
12 y3 = cp.Variable()
13
14 # Objective function for the primal problem
15 objective = cp.Maximize(1.5*x + 1.75*y + 2.0*z)
16
17 # Constraints for the primal problem
18 constraints = [
19     9*x + 4*y + z <= 1690,
20     5*x + 8*y + 5*z <= 1920,
21     2*x + 4*y + 10*z <= 1760,
22     x >= 0,
23     y >= 0,
24     z >= 0
25 ]
26
27 # Define the primal problem and solve it
28 primal_problem = cp.Problem(objective, constraints)
29 primal_problem.solve()
30
31 # Get dual variables from the primal constraints
32 dual_variables_from_primal = [constraint.dual_value for constraint in constraints]
33
34 # Objective function for the dual problem
35 dual_objective = cp.Minimize(1690*y1 + 1920*y2 + 1760*y3)
36
37 # Dual constraints (note that y1, y2, y3 should be non-negative)
```

```

38 dual_constraints = [
39     9*y1 + 5*y2 + 2*y3 >= 1.5,
40     4*y1 + 8*y2 + 4*y3 >= 1.75,
41     y1 + 5*y2 + 10*y3 >= 2,
42     y1 >= 0,
43     y2 >= 0,
44     y3 >= 0
45 ]
46
47 # Define the dual problem and solve it
48 dual_problem = cp.Problem(dual_objective, dual_constraints)
49 dual_problem.solve()
50
51 # Get dual variables from the dual constraints
52 primal_variables_from_dual = [constraint.dual_value for constraint in dual_constraints]
53
54 print("Primal Problem:")
55 print("Optimal value:", primal_problem.value)
56 print("Optimal solution:", (x.value, y.value, z.value))
57 print("Dual variables from primal constraints:", dual_variables_from_primal)
58
59 print("\nDual Problem:")
60 print("Optimal value:", dual_problem.value)
61 print("Optimal solution:", (y1.value, y2.value, y3.value))
62 print("Primal variables from dual constraints:", primal_variables_from_dual)
63
64 Primal Problem:
65 Optimal value: 581.5624999534823
66 Optimal solution: (array(139.06250001), array(80.62499996), array(115.9375))
67 Dual variables from primal constraints: [0.07291666670093609, 0.11458333334002731, 0.135416666683387, 2.111444868]
68
69 Dual Problem:
70 Optimal value: 581.5624999461611
71 Optimal solution: (array(0.07291667), array(0.11458333), array(0.13541667))
72 Primal variables from dual constraints: [139.06250001368736, 80.62499995648172, 115.93750000386534, 4.7030261270]

```

1.3 Objective value change

We expect our optimal value to increase. Here we are talking about resource limitations and by increasing the amount Brazil beans we have more of that resource. While the change of the objective value is hard to determine (without the code), just have more resources in our problem allows us for higher potential usage of beans/coffee which are beneficial for the objective.

The code (the only thing we change is 1600 to 1696) does reflect our instincts of the value increasing. Therefore in this optimization problem increasing the number of beans from 1600 to 1696 our objective value increases.

2 Problem 2

2.1 Model

We aim to Minimize the objective function:

$$3X_1 - X_2 + 2X_3$$

Subject to the constraints:

$$2X_1 + 3X_2 + 4X_3 \leq 10$$

$$X_1 + 5X_2 + X_3 = 4$$

$$-X_1 + 2X_2 + 6X_3 \geq 1$$

$$X_1, X_2, X_3 \geq 0,$$

To convert this problem into standard form, we must ensure all equations are strictly in the form of \leq some constant. The standard form is:

$$\min \quad 3X_1 - X_2 + 2X_3$$

subject to

$$2X_1 + 3X_2 + 4X_3 \leq 10,$$

$$X_1 + 5X_2 + X_3 \leq 4,$$

$$-X_1 - 5X_2 - X_3 \leq -4,$$

$$-X_1 + 2X_2 + 6X_3 \leq -1,$$

$$X_1, X_2, X_3 \geq 0.$$

Following the formula in the textbook, given by $\min / \max \mathbf{c}^T \lambda$ subject to $\mathbf{A}^T \lambda \geq \mathbf{b}, \lambda \geq 0$, and using α, β, γ , and δ as our λ vector, we can derive:

$$\min \quad 10\alpha + 4\beta - 4\gamma - \delta$$

subject to

$$2\alpha + \beta - \gamma + \delta \geq -3$$

$$-3\alpha - 5\beta + 5\gamma + 2\delta \geq -1$$

$$4\alpha + \beta - \gamma - 6\delta \geq -2$$

$$-4\alpha - \beta + \gamma + 6\delta \geq 2$$

$$\alpha, \beta, \gamma, \delta \geq 0.$$

2.2 Dual of the Dual

We can modify the signs in the objective function to convert from a minimization to a maximization problem. This provides the dual in the standard form which is typically a maximization problem.

$$\max \quad -10\alpha - 4\beta + 4\gamma + \delta$$

S.t:

$$\begin{aligned}
-2\alpha - \beta + \gamma - \delta &\leq 3, \\
3\alpha + 5\beta - 5\gamma - 2\delta &\leq 1, \\
-4\alpha - \beta + \gamma + 6\delta &\leq 2, \\
4\alpha + \beta - \gamma - 6\delta &\leq -2, \\
\alpha, \beta, \gamma, \delta &\geq 0
\end{aligned}$$

When we take dual of this we first switch this to a minimization problem

$$\min \quad -3y_1 + y_2 + 2y_3 - 2y_4$$

S.t:

$$\begin{aligned}
-2y_1 + 3y_2 - 4y_3 + 4y_4 &\geq -10 \\
-y_1 + 5y_2 - y_3 + y_4 &\geq -4 \\
y_1 - 5y_2 + y_3 - y_4 &\geq 4 \\
-y_1 - 2y_2 + 6y_3 - 6y_4 &\geq 1
\end{aligned}$$

When simplifying the model by combining y_3 and y_4 into a single variable x_3 and changing the inequality directions, we get:

$$\text{minimize} \quad 3x_1 + x_2 + 2x_3$$

Subject to:

$$\begin{aligned}
2x_1 + 3x_2 + 4x_3 &\leq 10 \\
x_1 + 5x_2 + x_3 &= 4 \\
-x_1 - 2x_2 + 6x_3 &\leq 1 \\
x_1 \geq 0, \quad x_2 \leq 0, \quad x_3 \text{ unrestricted}
\end{aligned}$$

2.3 Primal Optimal Solution

Given the optimal values of the dual variables:

$$\alpha^* = 0.5, \quad \beta^* = 0, \quad \gamma^* = 4, \quad \delta^* = 0$$

The constraints from the dual problem can be checked for activity:

$$\begin{aligned}
-2\alpha^* - \beta^* + \gamma^* - \delta^* &\leq 3 \\
3\alpha^* + 5\beta^* - 5\gamma^* - 2\delta^* &\leq 1 \\
-4\alpha^* - \beta^* + \gamma^* + 6\delta^* &\leq 2 \\
4\alpha^* + \beta^* - \gamma^* - 6\delta^* &\leq -2
\end{aligned}$$

Substituting the dual optimal values:

$$\begin{aligned}
-2(0.5) - 0 + 4 - 0 &= 3 \quad (\text{active}) \\
3(0.5) + 0 - 5(4) - 0 &= -18.5 \quad (\text{inactive}) \\
-4(0.5) - 0 + 4 + 0 &= 2 \quad (\text{active}) \\
4(0.5) + 0 - 4 - 0 &= -2 \quad (\text{active})
\end{aligned}$$

The corresponding primal constraints and their activity status determine the system of equations:

$$\begin{aligned} 2X_1 + 3X_2 + 4X_3 &= 10, \\ -X_1 + 2X_2 + 6X_3 &= 1, \\ X_1 + 5X_2 + X_3 &= 4. \end{aligned}$$

Using matrix operations or a system of equations solver, we can find the primal variables:

$$\begin{bmatrix} 2 & 3 & 4 \\ -1 & 2 & 6 \\ 1 & 5 & 1 \end{bmatrix} \begin{bmatrix} X_1 \\ X_2 \\ X_3 \end{bmatrix} = \begin{bmatrix} 10 \\ 1 \\ 4 \end{bmatrix}$$

From there we solve the system of equations to get $X_1 = 223/63$, $X_2 = -4/63$, and $X_3 = 7/9$

3 Problem 3

3.1 Dual of SVM

The primal problem for an SVM is given by:

$$\min \frac{1}{2} \|w\|^2 + C \sum_{i=1}^m e_i$$

subject to:

$$y_i(w^T x_i + b) \geq 1 - e_i, \quad i = 1, \dots, m$$

$$e_i \geq 0, \quad i = 1, \dots, m$$

Introduce dual variables $\alpha_i^{(1)}$ for the inequality $y_i(w^T x_i + b) \geq 1 - e_i$ and $\alpha_i^{(2)}$ for $e_i \geq 0$.

The Lagrangian for this problem is:

$$L(w, b, e, \alpha^{(1)}, \alpha^{(2)}) = \frac{1}{2} \|w\|^2 + C \sum_{i=1}^m e_i - \sum_{i=1}^m \alpha_i^{(1)} (y_i(w^T x_i + b) - 1 + e_i) - \sum_{i=1}^m \alpha_i^{(2)} e_i$$

Minimize L with respect to w, b, e by setting the derivatives to zero:

$$\frac{\partial L}{\partial w} = w - \sum_{i=1}^m \alpha_i^{(1)} y_i x_i = 0 \quad \Rightarrow \quad w = \sum_{i=1}^m \alpha_i^{(1)} y_i x_i$$

$$\frac{\partial L}{\partial b} = - \sum_{i=1}^m \alpha_i^{(1)} y_i = 0 \quad \Rightarrow \quad \sum_{i=1}^m \alpha_i^{(1)} y_i = 0$$

$$\frac{\partial L}{\partial e_i} = C - \alpha_i^{(1)} - \alpha_i^{(2)} = 0 \quad \Rightarrow \quad \alpha_i^{(1)} + \alpha_i^{(2)} = C$$

The dual problem then becomes:

$$\max_{\alpha^{(1)}} -\frac{1}{2} \sum_{i=1}^m \sum_{j=1}^m \alpha_i^{(1)} \alpha_j^{(1)} y_i y_j x_i^T x_j + \sum_{i=1}^m \alpha_i^{(1)}$$

subject to:

$$\sum_{i=1}^m \alpha_i^{(1)} y_i = 0, \quad 0 \leq \alpha_i^{(1)} \leq C, \quad i = 1, \dots, m$$

3.2 Optimal solution for w^*

$$w^* = \sum_{i=1}^m \alpha_i^* y_i x_i \tag{1}$$

where w^* is the sum of the vectors x_i , each weighted by their corresponding y_i and the dual variables α_i^* .

3.3 Extra Credit

extra credit

3.4 Dual model in CVXPY

```

1 import cvxpy as cp
2 import numpy as np
3 from sklearn import datasets
4 from sklearn.model_selection import train_test_split
5 from sklearn.preprocessing import StandardScaler
6
7 # Load dataset
8 data = datasets.load_breast_cancer()
9 X = data.data
10 y = data.target
11
12 # Convert labels from {0, 1} to {-1, 1}
13 y = 2 * y - 1
14
15 # Standardize data
16 scaler = StandardScaler()
17 X = scaler.fit_transform(X)
18
19 # Split data into training and test sets

```

```

20 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
21
22 # SVM parameters
23 C = 1.0
24 m = X_train.shape[0]
25
26 # Variables
27 w = cp.Variable(X_train.shape[1])
28 b = cp.Variable()
29 e = cp.Variable(m)
30
31 # SVM objective
32 objective = cp.Minimize(0.5 * cp.norm(w, 2) ** 2 + C * cp.sum(e))
33
34 # Constraints
35 constraints = [cp.multiply(y_train, X_train @ w + b) >= 1 - e, e >= 0]
36
37 # Define and solve the problem
38 problem = cp.Problem(objective, constraints)
39 problem.solve(solver=cp.ECOS)
40 # Output results
41 print("Status:", problem.status)
42 print("Optimal value:", problem.value)
43 print("w:", w.value)
44 print("b:", b.value)
45
46 # Optionally, calculate and print the accuracy on the test set if required
47 predictions = np.sign(X_test @ w.value + b.value)
48 accuracy = np.mean(predictions == y_test)
49 print("Test accuracy:", accuracy)
50 Status: optimal
51 Optimal value: 20.720270162367136
52 w: [-0.34852481 -0.04961735 -0.27412143 -0.22580064  0.04633631  0.74822593
53    -0.67740877 -1.41886349  0.1613721  -0.13479843 -0.9906163  0.38233107
54    -0.18813397 -0.78252804 -0.33523525  0.34831982  0.23050336 -0.17600259
55     0.40048013  0.59081354 -0.5963362  -1.21686748 -0.12895118 -0.5515148
56    -0.21737431  0.162141  -0.73861912 -0.13003616 -0.79205568 -0.53411574]
57 b: -0.0640986324235201
58 Test accuracy: 0.956140350877193

```


4 Problem 4

4.1 Update Rule for gradient Descent

Given the objective function:

$$f(x) = \frac{1}{2} \|Ax - b\|_2^2$$

where A is a matrix and b is a vector. We expand the objective function as:

$$f(x) = \frac{1}{2} (Ax - b)^T (Ax - b)$$

The gradient of the function (from class), $\nabla f(x)$, with respect to x is calculated as follows:

$$\nabla f(x) = A^T (Ax - b)$$

Thus, the update rule for gradient descent is:

$$x^{k+1} = x^k - \mu A^T (Ax^k - b)$$

If our derivative at some point $\|\Delta f(x_0)\|_2 < \epsilon$ then we are done (our answer converges), if not we keep iterating.

where:

- x^{k+1} is the value of x at iteration $k + 1$.
- x^k is the value of x at iteration k , which can also be the initial guess for $k = 0$.
- μ is the step size

4.2 Code

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3 from sklearn import datasets
4 from sklearn.preprocessing import StandardScaler
5
6 # Load the diabetes dataset
7 diabetes = datasets.load_diabetes()
8 X, y = diabetes.data, diabetes.target
9
10 # Standardize features
11 scaler = StandardScaler()
12 X = scaler.fit_transform(X)
13 y = y[:, np.newaxis] # Make y a column vector
```

```

14
15 # Append a column of ones to X for the intercept
16 X = np.hstack([np.ones((X.shape[0], 1)), X])
17
18 # Define the MSE loss function
19 def mse_loss(X, y, w):
20     m = len(y)
21     predictions = X @ w
22     loss = (1 / (2 * m)) * np.sum((predictions - y) ** 2)
23     gradient = (1 / m) * X.T @ (predictions - y)
24     return loss, gradient
25
26 # Gradient descent function that tracks error
27 def gradient_descent(X, y, step_size, max_iter=5000, tol=1e-3):
28     w = np.random.uniform(-0.1, 0.1, (X.shape[1], 1)) # Initialize weights
29     errors = []
30     for i in range(max_iter):
31         loss, grad = mse_loss(X, y, w)
32         errors.append(loss)
33         if np.linalg.norm(grad, 2) < tol:
34             break
35         w -= step_size * grad
36     return w, errors
37
38 # Test gradient descent with various step sizes
39 step_sizes = [1e-3, 1e-4, 1e-5, 1e-6]
40 all_errors = {}
41 for step in step_sizes:
42     _, errors = gradient_descent(X, y, step)
43     all_errors[step] = errors
44     print(f"Converged in {len(errors)} iterations for step size {step}")
45
46 # Run exact line search
47 exact_weights = line_search_gradient_descent(X, y)
48 print(f"Weights from exact line search:\n{exact_weights}")
49
50 # Plotting the errors
51 plt.figure(figsize=(10, 8))
52 for step, errors in all_errors.items():
53     plt.plot(errors, label=f'Step size {step}')
54
55 plt.xlabel('Iteration')
56 plt.ylabel('Training Error')
57 plt.title('Training Error vs. Iterations for Various Step Sizes')
58 plt.legend()
59 plt.yscale('log')
60 plt.grid(True)
61 plt.show()

```

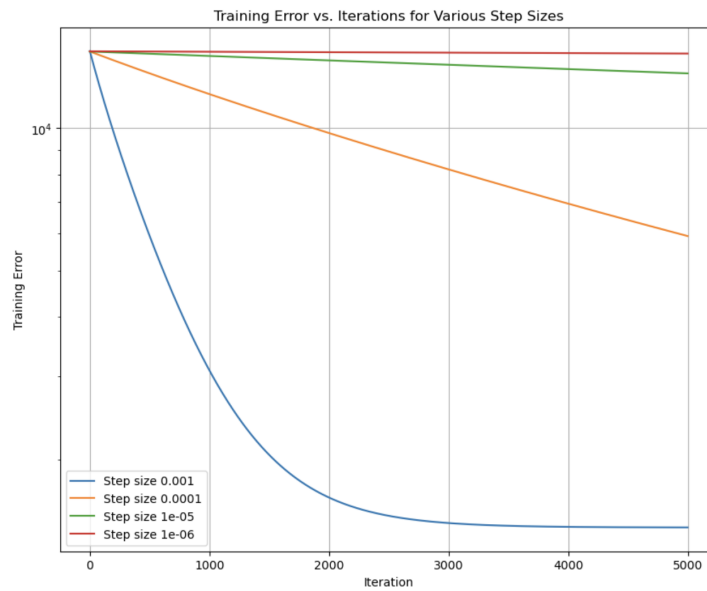


Figure 1: Graph 1

```

62 Converged in 5000 iterations for step size 0.001
63 Converged in 5000 iterations for step size 0.0001
64 Converged in 5000 iterations for step size 1e-05
65 Converged in 5000 iterations for step size 1e-06
66 Convergence achieved after 1475 iterations.
67 Weights from exact line search:
68 [[152.13348416]
69  [-0.47582498]
70  [-11.40654446]
71  [ 24.7272555 ]
72  [ 15.42919382]
73  [-37.62347913]
74  [ 22.63145722]
75  [  4.78085876]
76  [  8.41491072]
77  [ 35.71348882]
78  [  3.21693553]]

```

4.2.1 Explanation

In the context of gradient descent, the choice of step size, also known as the learning rate, plays a critical role in the convergence of the algorithm. When

the step size is set to a very small value, close to zero, the algorithm progresses towards the minimum of the loss function very slowly. This slow progress can result in the need for a significantly large number of iterations—potentially hundreds of thousands—to achieve a training error that is close to zero. This is because each update to the parameters of the model makes only a minimal adjustment, barely moving in the parameter space. On the other hand, when the step size is relatively larger, such as 0.001, the parameters are updated more substantially with each iteration. This more aggressive approach can substantially decrease the number of iterations needed to reach a satisfactory solution. For example, it might take only around 5000 to 6000 iterations to converge to a training error that is acceptably low. However, this approach is not without risks; too large a step size can lead to overshooting the minimum, causing the algorithm to diverge or oscillate around the minimum without settling. Thus, finding the right balance in step size is crucial. It should be large enough to ensure efficient convergence but not so large that it overshoots the minimum or fails to converge.

4.3 Comparison with CVXPY

```

1 import numpy as np
2 import cvxpy as cp
3 from sklearn import datasets
4 from sklearn.preprocessing import StandardScaler
5
6 # Load and normalize the diabetes dataset
7 diabetes = datasets.load_diabetes()
8 X = diabetes.data
9 y = diabetes.target
10
11 # Standardize features
12 scaler = StandardScaler()
13 X = scaler.fit_transform(X)
14 y = (y - np.mean(y)) / np.std(y) # Normalize target variable
15
16 # Append a column of ones to X for the intercept
17 X = np.hstack([np.ones((X.shape[0], 1)), X])
18
19 # Define and solve the CVXPY problem
20 n = X.shape[1]
21 x = cp.Variable(n)
22 cost = cp.norm(X @ x - y)
23 prob = cp.Problem(cp.Minimize(cost))
24 prob.solve(solver=cp.ECOS)
25
26 # Print results
27 print("\nThe optimal value is", prob.value)

```

```

28 print("The optimal x is")
29 print(x.value)
30 print("The norm of the residual is ", cp.norm(X @ x - y, p=2).value)
31 CVXPY SVM Status: optimal
32 CVXPY SVM Optimal value: 20.720270162367136
33 CVXPY SVM w: [-0.34852481 -0.04961735 -0.27412143 -0.22580064  0.04633631  0.74822593
34 -0.67740877 -1.41886349  0.1613721  -0.13479843 -0.9906163  0.38233107
35 -0.18813397 -0.78252804 -0.33523525  0.34831982  0.23050336 -0.17600259
36  0.40048013  0.59081354 -0.5963362  -1.21686748 -0.12895118 -0.5515148
37 -0.21737431  0.162141  -0.73861912 -0.13003616 -0.79205568 -0.53411574]
38 CVXPY SVM b: -0.0640986324235201
39 CVXPY SVM Test accuracy: 0.956140350877193

```

4.3.1 Explanation

We see that these results don't exactly match the results from 2B and the major reason why they don't match is because we are using two approaches for optimizing data. Our current SVM model is trying to classify data points as either having cancer or not, using a hyperplane defined by w and b . While our

linear regression/least squares approximation is trying to minimize the residuals of the predicted value and actual values of people with cancer.

Additionally, our constraints are vastly different too. We see classification boundaries and slack for SVM vs usually unconstrained for simple regression.

4.4 d

extra credit

Appendix

I used AI to debug some of coding issues